

# Lab 1A Handout 2024

## Starter Lab

### Lab Overview

Lab 1A is the starter lab, designed to introduce the development tools used for creating hardware/software interfaces. After completing Lab 1A, you will have created an RTL project using the **Vivado** design suite from Xilinx. Your basic RTL project will consist of the soft-core **Microblaze**[2] processor running on the **Artix A7-100T development board**[1] from Digilent. The Artix A7 board utilizes a Xilinx FPGA chip, the Artix 7. The Artix A7 board was previously named Nexys 4 DDR, and is identical to the Nexys 4 DDR in terms of functionality.

The Microblaze IP core is configurable, and permits creation of a system with different memory sizes, clock speeds and # of pipeline stages. The goal of this lab is to use the Vivado design Suite to create a custom hardware design for a processor and its peripherals, export the hardware to Xilinx Vitis, and then develop **C** programs to run on the custom processor. The **C** programs will gather timing information about the embedded system in terms of the number of clock cycles to execute different operations. Finally, timing uncertainties in the gathered data will be statistically analyzed.

## 1 Equipment

Throughout the course we will be creating Vivado RT (Register Transfer) projects, which are based around the block level diagram. An RT project refers to the level of abstraction selected for describing the design.

**Hardware** Nexys A7-100T Digilent Development board with Artix-7 FPGA

**Software** Vivado 2022.1

**Software** Xilinx Vitis (with Vivado Integration)

The **Lab Hardware Handout** provides information about setting up the Artix A7 Development board with Vivado 2022.1 and the Xilinx Vitis. It also contains information on interfacing the DDR2 memory with Microblaze, and general information which is useful for all of the labs. For this lab, the DDR2 memory will need to be a part of your design. Read the **Lab Hardware Handout** before starting the exercises. You can control how much and which memory is used (either BRAM or DDR2) by modifying the linker file created as part of the software design. Doing so has consequences on the throughput and latency of memory access for your software, thus directly changing the results of your testing. Please pay particular attention to the settings and enabled routines when collecting your data.

## 2 Microblaze Configuration to be Explored

The configuration of the Microblaze processor provides various options, such as the size of the BRAM memory used by Microblaze, adding data and instruction caches and enabling debugging tools. For this lab, build Microblaze using the

configuration instructions in the Lab hardware handout. Follow instructions in the **Lab Hardware Handout** to start Vivado and create a project containing a Microblaze processor along with DDR2 Memory. For this lab you should run your program in the DDR memory with caches enabled. Note: when re-running software for debugging or otherwise, be particularly careful you invalidate any cache before you enable it. If you don't do this, the program will assume that the current cache values are correct – and they might well conflict with updated program test or code changes leading to unpredictable behavior or instability.

### 3 Adding Peripherals

The Microblaze processor provides a central processing unit for our embedded system to interact with the physical world. The next step in the design of the system is to begin adding I/O peripherals. Follow instructions in the **Lab Hardware Handout** to add IP for: (1) Two AXI Timers and (2) LEDs. You should add the other peripherals in the hardware handout, but they are not used in this lab. Peripherals communicate with the processor using buses. In our design we are using the AXI bus. Buses are one of the design elements which face timing issues. As we complete the timing analysis of operations listed in Section 5, we begin to see when peripherals are competing for bus access.

### 4 Timing Analysis Guidelines

After building the hardware and generating bitstream in Vivado, export your hardware and launch Vitis. Create a new Application Project in Vitis. We have provided default source code files that you will need for this lab. Copy these files to the “src” directory of your Application Project. **\*\*\*Brandon, please look at this bit and ensure the instructions are updated!\*\*\*** The main code for this project that will be edited by you is enclosed in timing.c. This code provides you with the basic structure for exercising different peripherals and for measuring the number of clock cycles.

Timing measurement on a computer is rather more complex than you might think. Abstractly, we can just read the timer, then do something and then read the new timer and voila... In real life, the memory access to the data and the instruction access are both non-zero time behaviors, the call to read the timer involves a context switch and stack update... In short, there is a lot of random and some systematic overhead. In particular, timing a single operation or event is hard because you have substantial random measurement overhead and most operations are designed to run faster if repeated (this is the fundamental premise of caching). One solution is to in-line write 5 events, 10 events, 20 events as so on – then do a regression to estimate the actual time it takes for events to run. Linear regression builds a ‘fit’ of the measurement data and a linear model based on two parameters to best model the data. The two parameters can be thought of as ‘overhead’ (the constant delay any measurement takes) and ‘delay per operation’ (the amount of additional time taken as the number of operations is increased). Such a model make the presumption that the measured quantity is indeed linear, something that should be verified by plotting the measurement and the fit.

A variable ‘Data’ has been setup in the default code to run such analysis on DDR read operations. Change the number of in-line calls in the code and fit a line (use linear regression) for 5, 15, 25, 40 in-line calls to estimate the measurement overhead and the per read timing of DDR memory. DDR timing is complicated by the caching involved – you will get substantially different answers depending on the addresses you access and where they are sequentially. You can now use this estimation of average run time to offset your test measurements. This way you are estimating the actual run time of the function you want to test. Note the noise associated with this measurement. Do you see what you expect? Note, when setting up the debug

configuration, you can set menay things– you can ask that the firware bitstream be re-loaded on reset (or not– it does not need reloading if you are only changing software an saves time on non-local reloads), you can also tell the process to stop on `main()` or not – and finally, you can setup a console logging file to store all data reported by the card when run, and where to put it. This saves a lot of manual copying if you are dumping histogram data from the card.

For the list below, measure each operation repeatedly using the built-in loop. Do not inline the timed operations, use the measurement overhead above to estimate the actual event times. Run the loop at least 10000 times to get a stab at some of the more rare event timings, for the USB port, run the loop 1000 times to keep the experiment time reasonable. **Be careful not to simply increase the loop too much as the default data storage to the histogram can overwhelm the stack.** This will cause unpredictable errors in run-time. (Of course, you can freely expand the stack in the linker file...).

A histogram function that counts the number of samples in uniformly spaced, equally sized bins is included in the code. You may choose an appropriate value for total number of bins and use this function for statistically analyzing your data. Be sure to use enough bins to display the variance in the data. You can also modify the code to change bin spacing or data usage and printout. (Marking the printout makes the extraction of data from the log file easy.) Finally, it is a good idea to analyze and plot your results using computer resources. I recommend Python or Mathematica™ as tools (python is free and has extensive package support - and runs on just about anything), Mathematica is supported by a site-wide license at UCSB, and runs in the lab or remotely given your license. They are so much nicer than data tasks by hand.

## 5 List of Operations to be Timed

1. Timing of addition using integers
2. Timing of addition using floating point
3. Timing of writing the LEDs to turn on or off
4. Timing of reading a word from the DDR2 memory at a random location
5. Timing of writing to the USB Port: (**You need to run the loop for this part only 1000 times**)
  - (a) Write a floating point number using *printf()* (Always use `printf()` with a “\n”, eg.: `printf(“A\n”)` to avoid BRAM overflow. `fprintf` also can work if “\n” is not working. The issue here is that `printf()` is a buffered I/O call, so you can keep adding to the buffer (which is usually on the stack) by not adding the return which usually acts as a signal to write the buffer.
  - (b) Write a string of 10 characters using *xil\_printf()*

While you are gathering data, another method in `main`, `extra_method()` is also running. This `extra_method()` adds a source of resource contention, which will change the time for communicating with these peripherals.

## 6 Reporting Results and Observations

It should be taken into consideration that embedded systems are targeted at tightly constrained environments, with multi-kHz sampling rates, multiple interrupt sources, and relatively small memories with realistic bus environments. Write a two-page report (note you can add an additional page for plots, but use space wisely), including the following results in your report:

1. Report the regression estimate of overhead and DDR timing from section 4.

2. Include histograms of the measurement data for the operations listed in Section 5.
3. Plot a CDF for the DDR access timing, What is the expected value for this timing? What is a 95% confidence interval for this measurement?
4. What do you observe about the magnitude of measured values and the timing uncertainties associated with different operations?
5. Did you have any outliers in your results? If so, what could have caused these values?
6. Can you explain why certain operations show smaller variations as compared to others?

## References

- [1] *Artix A7-100T FPGA Board*. URL: <https://store.digilentinc.com/nexys-a7-fpga-trainer-board-recommended-for-ece-curriculum/>.
- [2] *ECE 153A - Course Microblaze Notes*. URL: [Canvas::MicroBlaze\\_Overview.pdf](#).